OS security - OS key concepts

Myrto Arapinis
School of Informatics
University of Edinburgh

March 6, 2019

Operating systems

- An OS provides the interface between the users of a computer and that computer's hardware.
- The OS handles the management of low-level hardware resources:
  - disk drives,
  - CPU,
  - RAM,
  - I/O devices, and

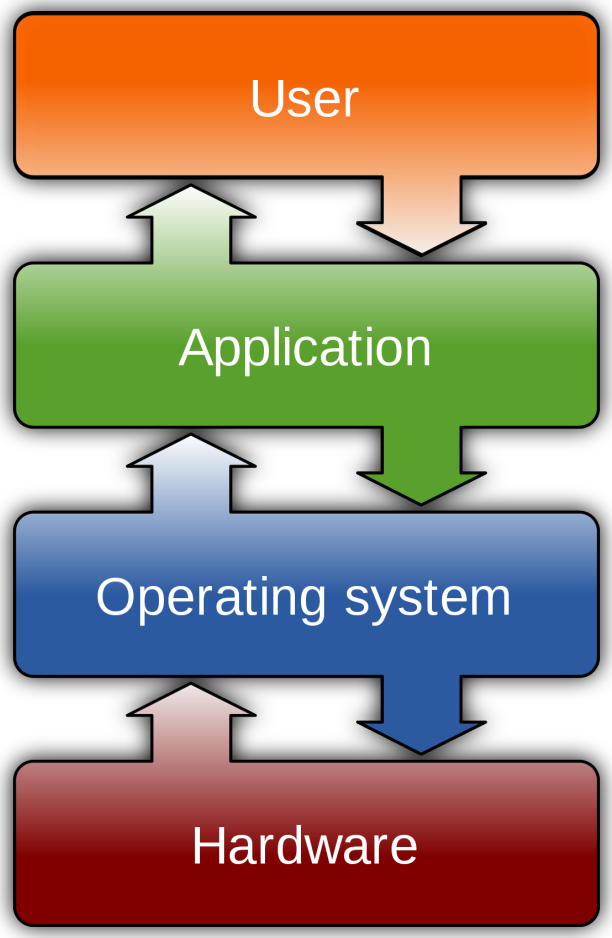

Multi-users

OSs must allow for multiple users with potentially different levels of access to the same computer.

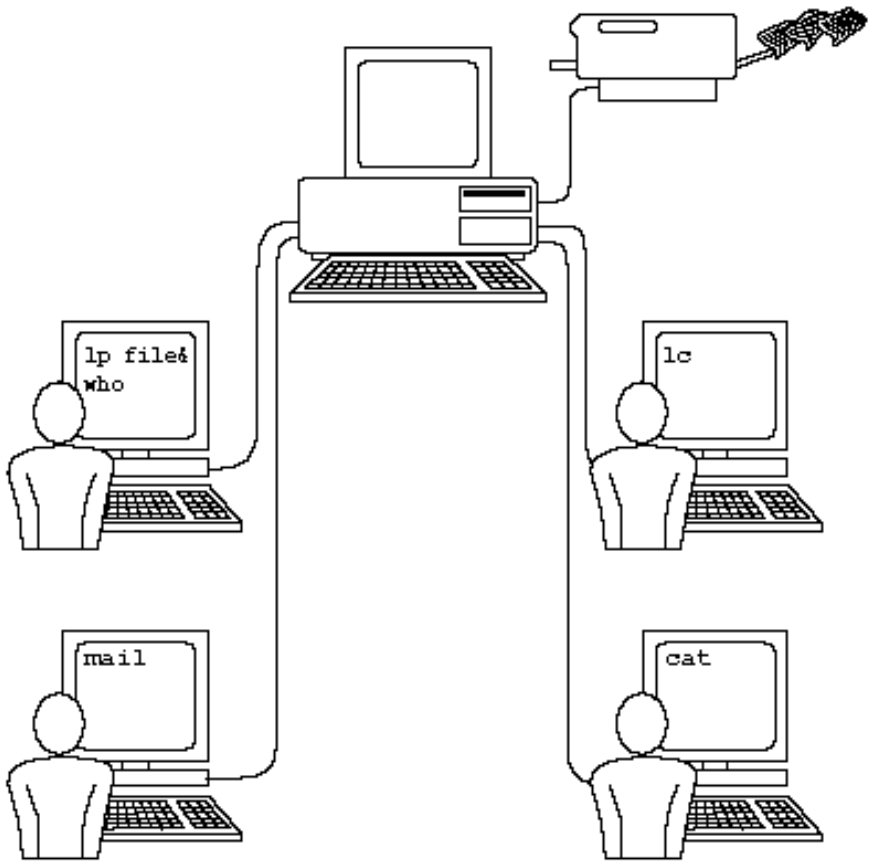

Multi-tasking

OSs must allow multiple application programs to run at the same time.

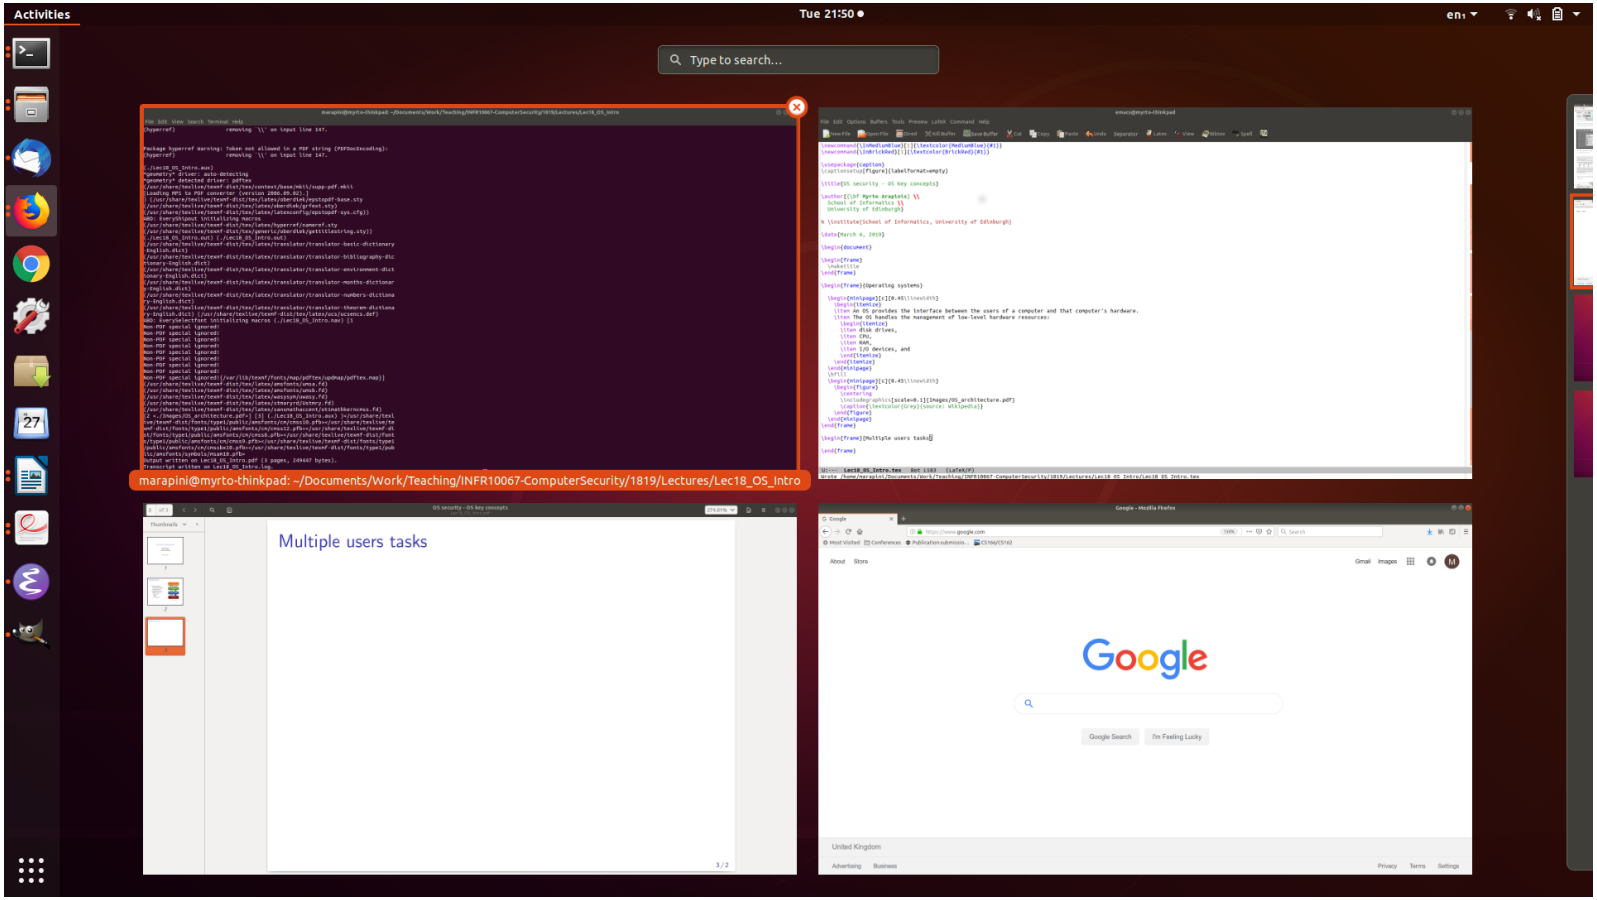

Essential Unix architecture

Execution modes:

- User mode - access to resources through syscall to kernel
- Kernel mode - direct access to resources

System calls are usually contained in a collection of programs, eg. a library such as the C library `libc`

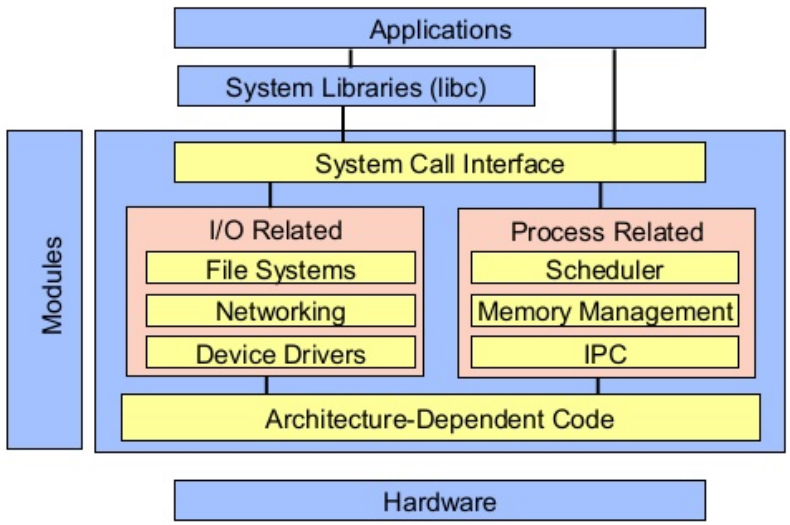

Processes and process management

- A process is an instance of a program that is currently executing.
- To actually be executed the program must be loaded into RAM and uniquely identified.
- Each process running is identified by a unique Process ID (PID).
- To a PID, we can associate its CPU time, memory usage, user ID (UID), program name, etc.
- A process might control other processes (fork).
- Child process inherits context from parent process.

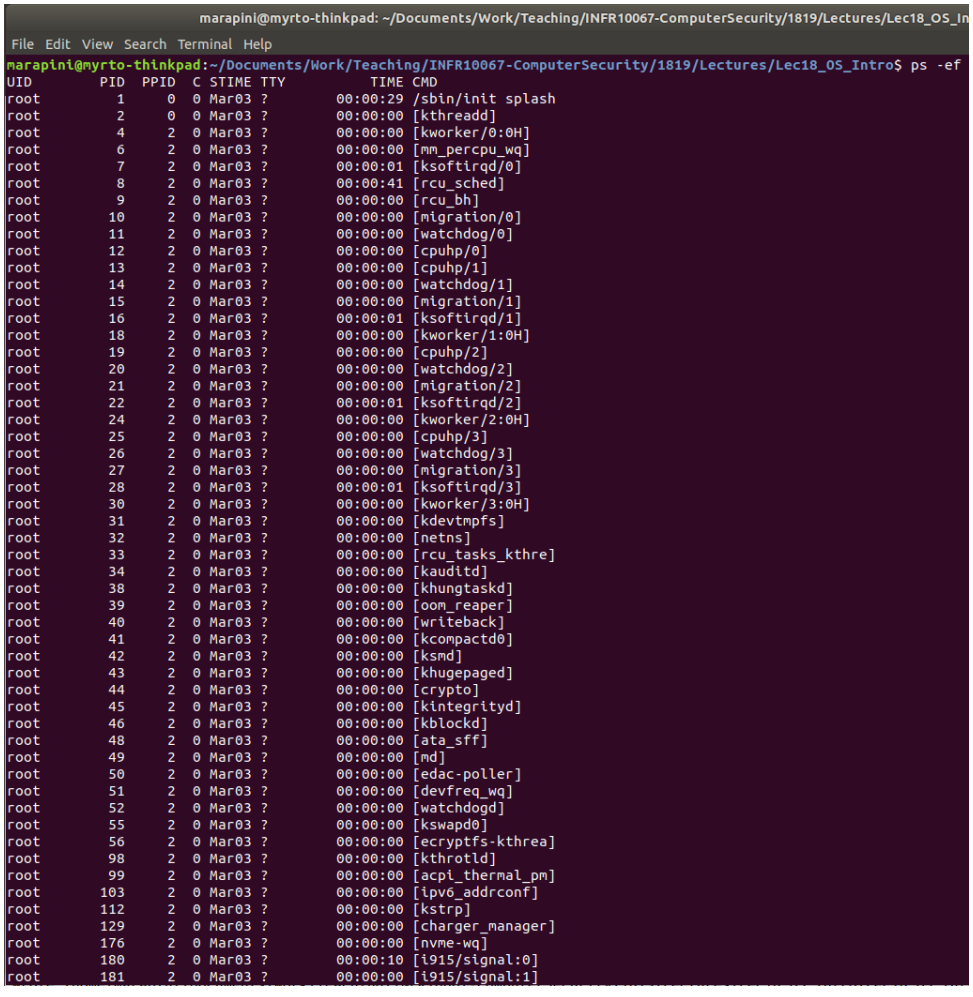

File permissions

- One of the main concern of OSs is how to delineate which user can access which resources.
- File permissions are checked by the OS to determine if a file is readable, writable, or executable by a user or a group of users.

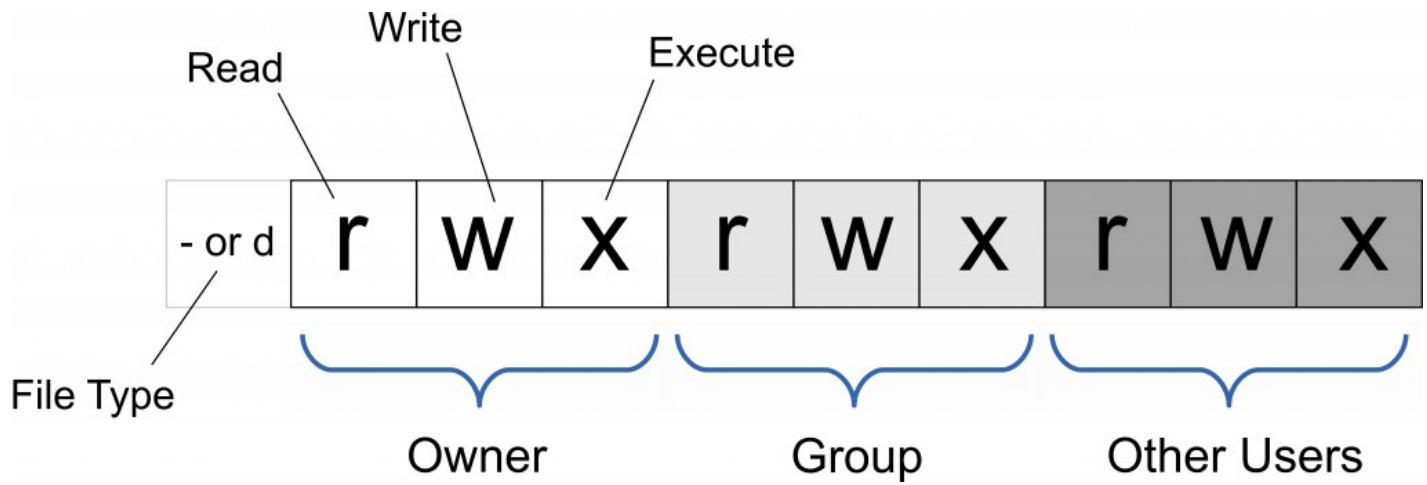

Unix file permissions

Setuid programs

- Unix process have 2 user IDs:
  - real user ID (uid) - user launching the pgm
  - effective user ID (euid) - user owning the pgm
- An executable file can have the set-user-ID property (setuid) enabled
- If A executes setuid file owned by B, then the euid of the process is B and not A
- Syscall setuid(uid) allows a process to change its euid to uid
- Writing secure setuid programs is tricky because vulnerabilities may be exploited by malicious user actions
- Some programs that access system resources are owned by root and have the setuid bit set (setuid programs)

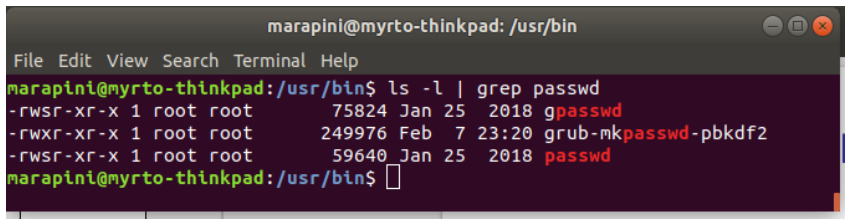

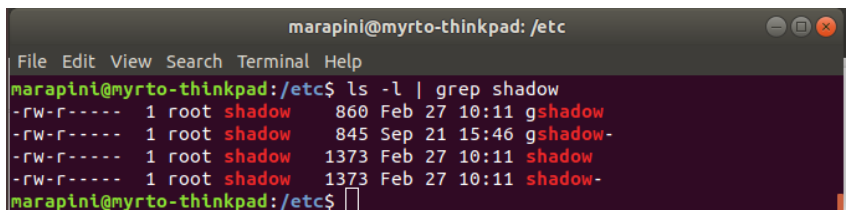

Memory management

- To actually be executed the program must be loaded into RAM and uniquely identified.
- The RAM memory of a computer is its address space.
- It contains both the code for the running program, its input data, and its working memory.
- For any running process it is organised into different segments, which keep the different parts of the address space separate
- Security concerns require that we do not mix up these different segments.

Linux (32-bit) process memory layout (simplified)

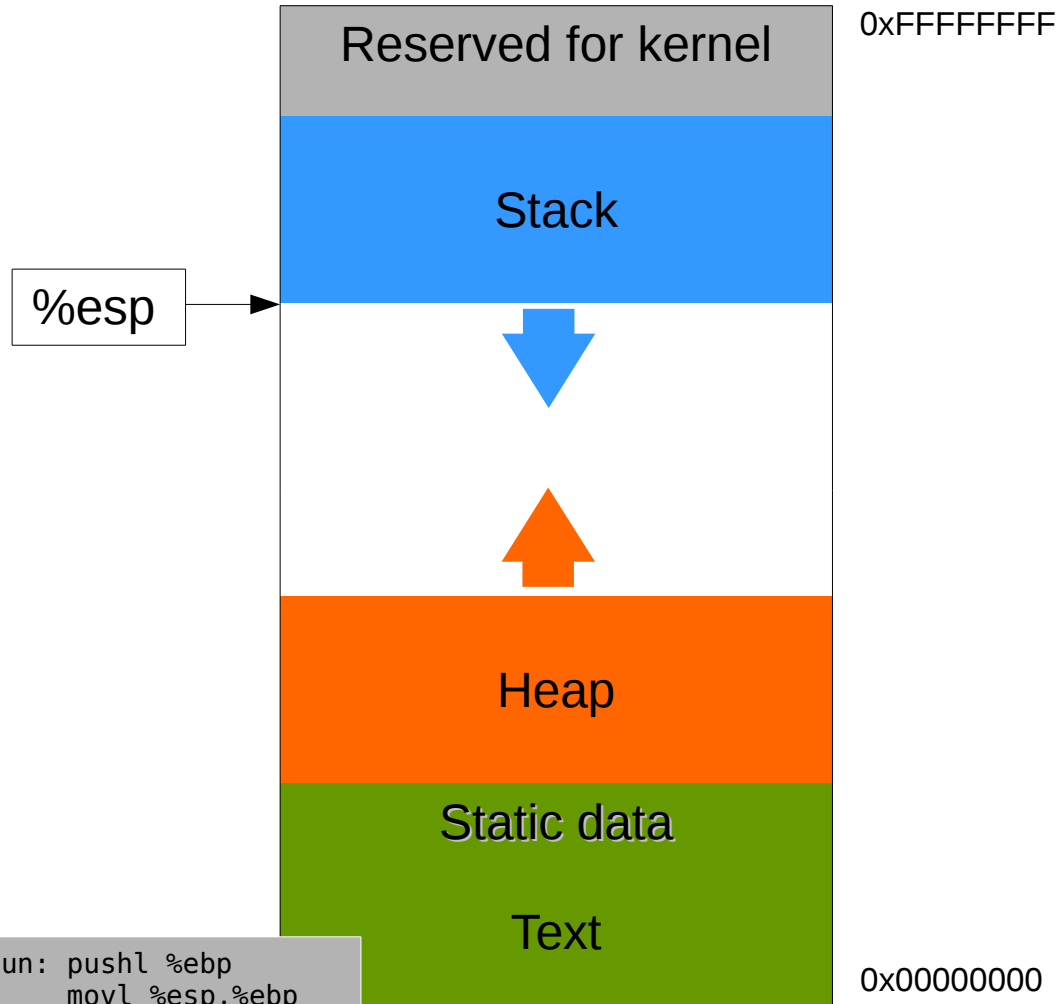


```
void function(int a, int b, int c) {
     char buffer1[5];
     char buffer2[10];
}
void main() {
     function(1,2,3);
}
```

```
fun: pushl %ebp
     movl %esp,%ebp
     subl $20,%esp

main: pushl $3
      pushl $2
      pushl $1
      call fun
```

Stack frame

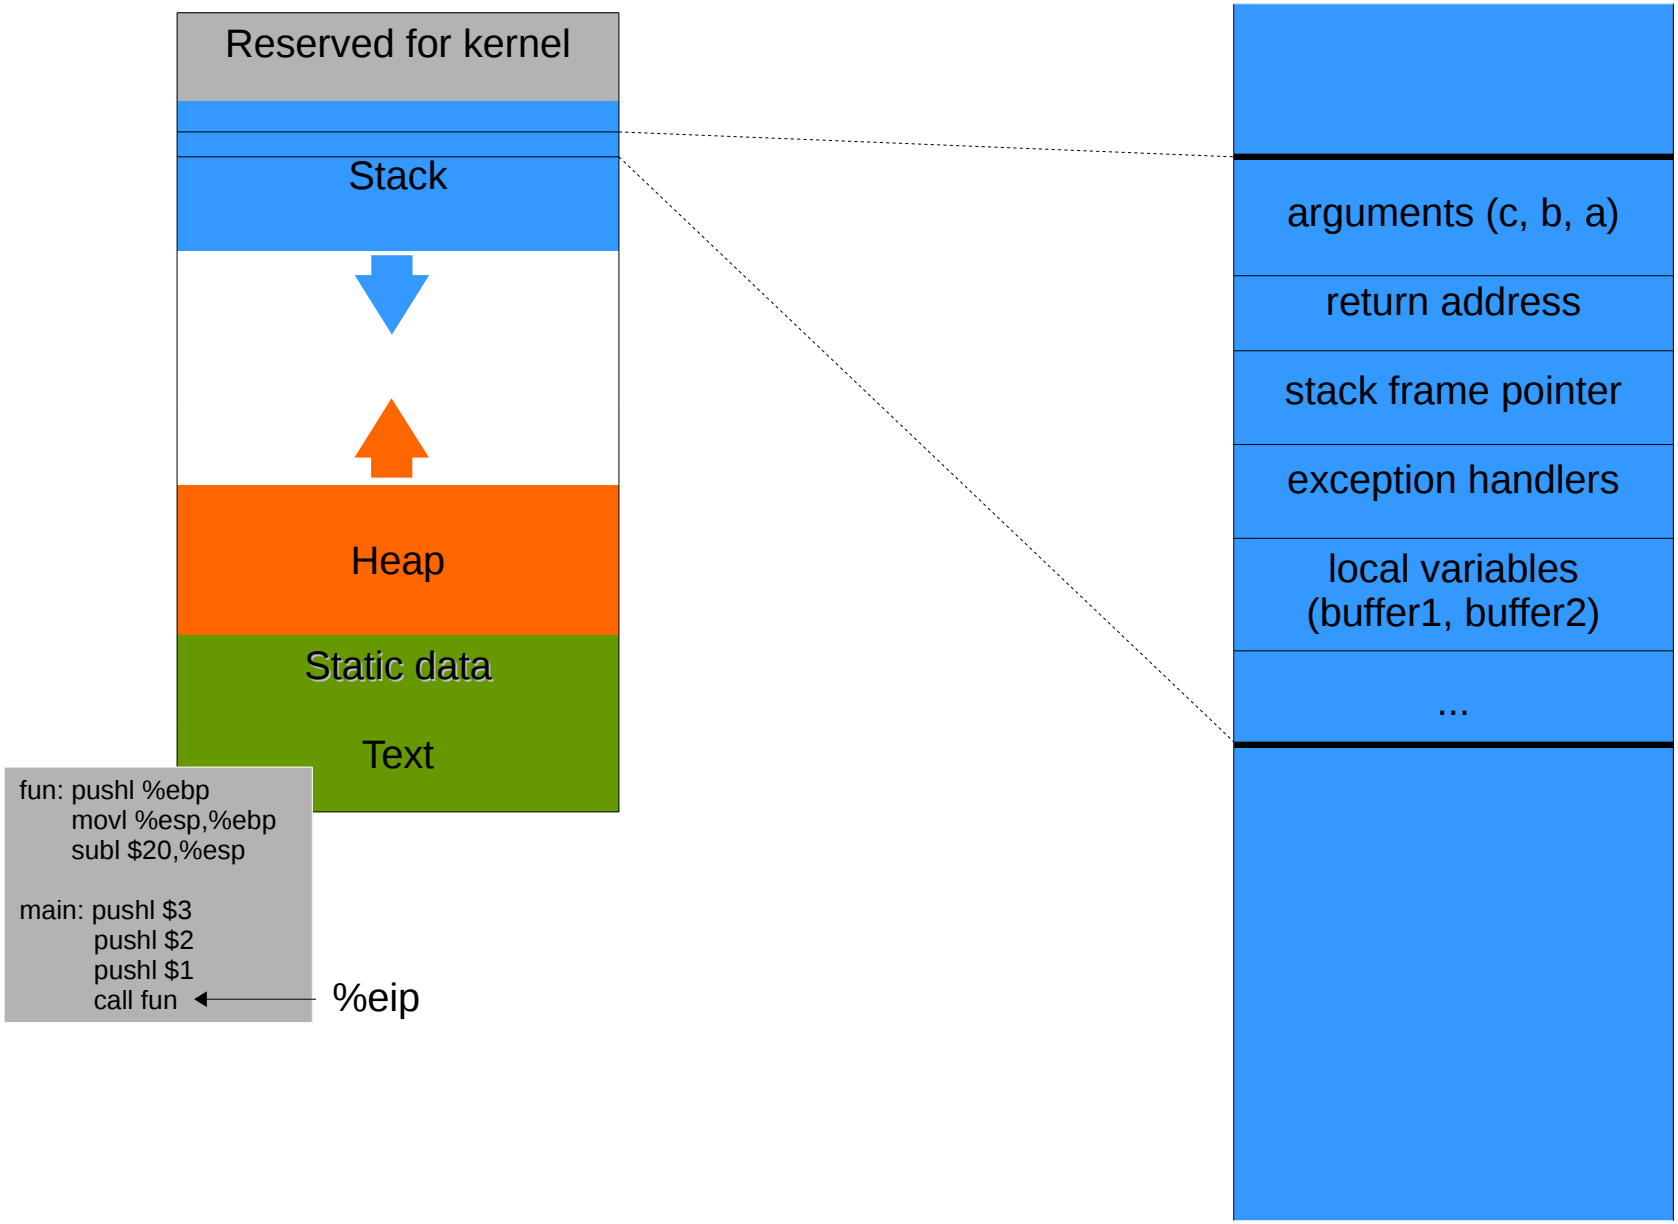
Reserved for kernel
Stack
Heap
Static data
Text
fun: pushl %ebp
movl %esp,%ebp
subl $20,%esp
main: pushl $3
pushl $2
pushl $1
call fun
%eip
arguments (c, b, a)
return address
stack frame pointer
exception handlers
local variables
(buffer1, buffer2)
...

Stack and functions: Summary

Calling function

1. Push arguments onto the stack (in reverse)
2. Push the return address, i.e., the address of the instruction to run after control returns
3. Jump to the function's address

Called function

4. Push the old frame pointer onto the stack (%ebp)
5. Set frame pointer (%ebp) to where the end of the stack is right now (%esp)
6. Push local variables onto the stack

Returning function

7. Reset the previous stack frame: %esp = %ebp, %ebp = (%ebp)
8. Jump back to return address: %eip = 4(%esp)