


Buffer overflows

Myrto Arapinis
School of Informatics
University of Edinburgh

March 08, 2019

x86 CPU/Memory

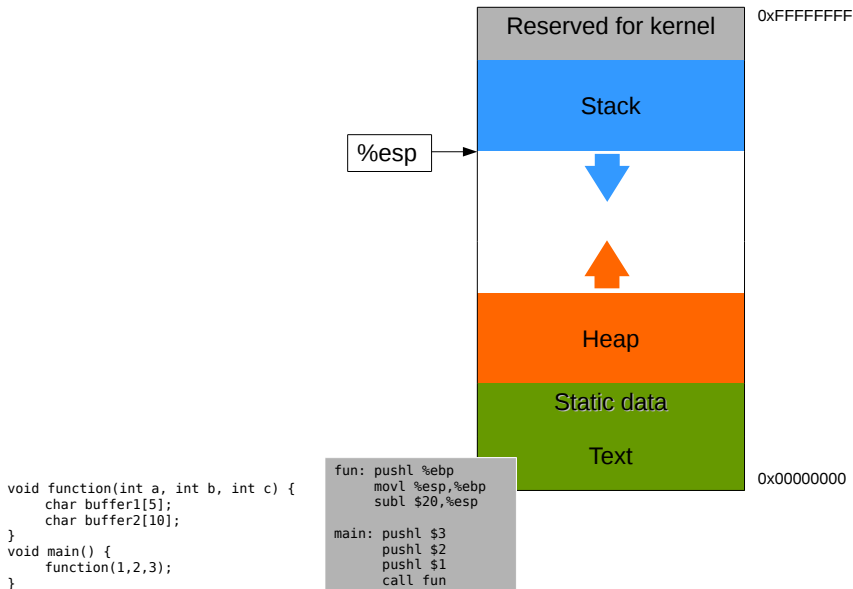


- ▶ Memory stores instructions and data
- ▶ CPU interprets instructions
- ▶ `%eip` points to next instruction
- ▶ `%eip` incremented after each instruction
- ▶ `%eip` modified by `call`, `ret`, `jmp`, and conditional `jmp`

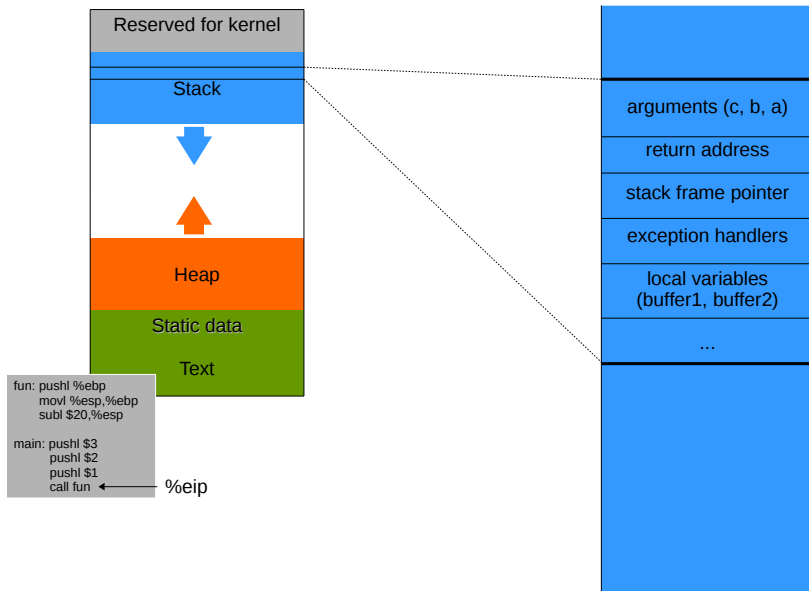
x86 registers

- ▶ Temporary registers: `%eax`, `%ebx`, `%ecx`, `%edx`, `%edi`, `%esi`
- ▶ Extended stack pointer: `%esp`
- ▶ Extended base pointer: `%ebp`

x86 process memory layout (simplified)



Stack frame



Stack and functions: Summary

Calling function

1. Push arguments onto the stack (in reverse)
2. Push the return address, i.e., the address of the instruction to run after control returns
3. Jump to the function's address

Called function

4. Push the old frame pointer onto the stack (`%ebp`)
5. Set frame pointer (`%ebp`) to where the end of the stack is right now (`%esp`)
6. Push local variables onto the stack

Returning function

7. Reset the previous stack frame: `%esp = %ebp, %ebp = (%ebp)`
8. Jump back to return address: `%eip = 4(%esp)`

x86 assembly

```
emac@myrto-thinkpad
File Edit Options Buffers Tools C Help

#include <stdio.h>

int fact(int x) {
    int y = 1;
    if (x == 0)
        return y;
    if (x > 0) {
        y = fact(x-1);
        return x*y;
    }
}

int main(void){
    int x = 0;
    x = fact(4);
    printf("Factorial 4 is %d\n", x);
    return 0;
}

--:--- Fact.c      All L1      (C/I Abbrev)
```

```
emac@myrto-thinkpad
File Edit Options Buffers Tools Asm Help

fact:
    pushl   %ebp
    movl    %esp, %ebp
    subl    $40, %esp
    movl    $1, -12(%ebp)
    cmpl    $0, 8(%ebp)
    jne     .L2
    movl    -12(%ebp), %eax
    jmp     .L1

.L2:
    jle     .L1
    movl    8(%ebp), %eax
    subl    $1, %eax
    movl    %eax, (%esp)
    call    fact
    movl    %eax, -12(%ebp)
    movl    8(%ebp), %eax
    null    -12(%ebp), %eax
    jmp     .L1

.L1:
    leave
    ret


.LC0:
    .string "Factorial 4 is %d\n"

main:
    pushl   %ebp
    movl    %esp, %ebp
    andl    $-16, %esp
    subl    $32, %esp
    movl    $0, 28(%esp)
    movl    $4, (%esp)
    call    fact
    movl    %eax, 28(%esp)
    movl    28(%esp), %eax
    movl    %eax, 4(%esp)
    movl    $.LC0, (%esp)
    call    printf
    movl    $0, %eax
    leave
    ret

--:--- Fact.c      All L26      (Assembler)
Write /home/naraptnl/Documents/Work/Teaching/INFR10067-ComputerSecurity/1819/Lecture/Lec18_05_intro/CDB_demo/fact.s
```

x86 runtime memory

```
int fact(int x) {  
    int y = 1;  
    if (x == 0)  
        return y;  
    if (x > 0) {  
        y = fact(x-1);  
        return x*y;  
    }  
}  
  
int main(void) {  
    int x = 0;  
    x = fact(4);  
    printf("Factorial 4 is %d", x);  
    return 0;  
}
```



x86 runtime memory

```
int fact(int x) {  
    int y = 1;  
    if (x == 0)  
        return y;  
    if (x > 0) {  
        y = fact(x-1);  
        return x*y;  
    }  
}  
  
int main(void) {  
    int x = 0;  
    x = fact(4);  
    printf("Factorial 4 is %d", x);  
    return 0;  
}
```

		...
<i>SF_{main}</i> —	0xbffeffc : 0xbffeff8 : 0xbffeff4 :	0xb7e31a83 @ret0 0x00000000 %ebp0 0x00000000

x86 runtime memory

```
int fact(int x) {  
    int y = 1;  
    if (x == 0)  
        return y;  
    if (x > 0) {  
        y = fact(x-1);  
        return x*y;  
    }  
}  
  
int main(void) {  
    int x = 0;  
    x = fact(4);  
    printf("Factorial 4 is %d", x);  
    return 0;  
}
```

		...
<i>SF_{main}</i> —	0xbffeffc : 0xbffeff8 : 0xbffeff4 :	0xb7e31a83 @ret0 0x00000000 %ebp0 0x00000000
<i>SF_{fact(4)}</i> —		

x86 runtime memory

```
int fact(int x) {
    int y = 1;
    if (x == 0)
        return y;
    if (x > 0) {
        y = fact(x-1);
        return x*y;
    }
}

int main(void) {
    int x = 0;
    x = fact(4);
    printf("Factorial 4 is %d", x);
    return 0;
}
```

		...
$SF_{main} -$	0xbffffeffc : 0xbffffeff8 : 0xbffffeff4 :	0xb7e31a83 @ret0 0x00000000 %ebp0 0x00000000
$SF_{fact(4)} -$	0xbffffefd0 : 0xbffffefcc : 0xbffffefc8 : 0xbffffefbc :	0x00000004 0x08048474 @retm 0xbffffeff8 %ebp_m 0x00000001

x86 runtime memory

```
int fact(int x) {
    int y = 1;
    if (x == 0)
        return y;
    if (x > 0) {
        y = fact(x-1);
        return x*y;
    }
}

int main(void) {
    int x = 0;
    x = fact(4);
    printf("Factorial 4 is %d", x);
    return 0;
}
```

		...
SF_{main} —	0xbffffc : 0xb7e31a83 @ret0 0xbffff8 : 0x00000000 %ebp0 0xbffff4 : 0x00000000	
$SF_{fact(4)}$ —	0xbffefd0 : 0x00000004 0xbffefcc : 0x08048474 @retm 0xbffefc8 : 0xbffff8 %ebp0 0xbffefbc : 0x00000001	
$SF_{fact(3)}$ —		

x86 runtime memory

```
int fact(int x) {
    int y = 1;
    if (x == 0)
        return y;
    if (x > 0) {
        y = fact(x-1);
        return x*y;
    }
}

int main(void) {
    int x = 0;
    x = fact(4);
    printf("Factorial 4 is %d", x);
    return 0;
}
```

		...
SF_{main} —	0xbffffc : 0xb7e31a83 @ret0 0xbffff8 : 0x00000000 %ebp0 0xbffff4 : 0x00000000	
$SF_{fact(4)}$ —	0xbffffd0 : 0x00000004 0xbffffc : 0x08048474 @retm 0xbffff8 : 0xbffff8 %ebp _m 0xbffffbc : 0x00000001	
$SF_{fact(3)}$ —	0xbffffa0 : 0x00000003 0xbffff9c : 0x08048449 @ret ₄ 0xbffff98 : 0xbffff8 %ebp ₄ 0xbffff8c : 0x00000001	

x86 runtime memory

```
int fact(int x) {
    int y = 1;
    if (x == 0)
        return y;
    if (x > 0) {
        y = fact(x-1);
        return x*y;
    }
}

int main(void) {
    int x = 0;
    x = fact(4);
    printf("Factorial 4 is %d", x);
    return 0;
}
```

		...
SF_{main} —	0xbfffffc : 0xb7e31a83 @ret ₀ 0xbfffff8 : 0x00000000 %ebp ₀ 0xbfffff4 : 0x00000000	
$SF_{fact(4)}$ —	0xbffffd0 : 0x00000004 0xbffffcc : 0x08048474 @ret _m 0xbffffc8 : 0xbfffff8 %ebp _m 0xbffffbc : 0x00000001	
$SF_{fact(3)}$ —	0xbffffa0 : 0x00000003 0xbffff9c : 0x08048449 @ret ₄ 0xbffff98 : 0xbffffc8 %ebp ₄ 0xbffff8c : 0x00000001	
$SF_{fact(2)}$ —		

x86 runtime memory

```
int fact(int x) {
    int y = 1;
    if (x == 0)
        return y;
    if (x > 0) {
        y = fact(x-1);
        return x*y;
    }
}

int main(void) {
    int x = 0;
    x = fact(4);
    printf("Factorial 4 is %d", x);
    return 0;
}
```

		...
SF_{main} —	0xbffffc : 0xb7e31a83 @ret0 0xbffff8 : 0x00000000 %ebp0 0xbffff4 : 0x00000000	
$SF_{fact(4)}$ —	0xbfffd0 : 0x00000004 0xbfffc : 0x08048474 @retm 0xbfffc8 : 0xbffff8 %ebp _m 0xbfffb : 0x00000001	
$SF_{fact(3)}$ —	0xbfffa0 : 0x00000003 0xbff9c : 0x08048449 @ret ₄ 0xbff98 : 0xbfffc8 %ebp ₄ 0xbff8c : 0x00000001	
$SF_{fact(2)}$ —	0xbff70 : 0x00000002 0xbff6c : 0x08048449 @ret ₃ 0xbff68 : 0xbff98 %ebp ₃ 0xbff5c : 0x00000001	

x86 runtime memory

```
int fact(int x) {
    int y = 1;
    if (x == 0)
        return y;
    if (x > 0) {
        y = fact(x-1);
        return x*y;
    }
}

int main(void) {
    int x = 0;
    x = fact(4);
    printf("Factorial 4 is %d", x);
    return 0;
}
```

		...
SF_{main} —	0xbffffc : 0xb7e31a83 @ret0 0xbffff8 : 0x00000000 %ebp0 0xbffff4 : 0x00000000	
$SF_{fact(4)}$ —	0xbfffd0 : 0x00000004 0xbffefc : 0x08048474 @retm 0xbffefc8 : 0xbffff8 %ebp_m 0xbffefbc : 0x00000001	
$SF_{fact(3)}$ —	0xbffefa0 : 0x00000003 0xbffef9c : 0x08048449 @ret4 0xbffef98 : 0xbffefc8 %ebp4 0xbffef8c : 0x00000001	
$SF_{fact(2)}$ —	0xbffef70 : 0x00000002 0xbffef6c : 0x08048449 @ret3 0xbffef68 : 0xbffef98 %ebp3 0xbffef5c : 0x00000001	
$SF_{fact(1)}$ —		

x86 runtime memory

```
int fact(int x) {
    int y = 1;
    if (x == 0)
        return y;
    if (x > 0) {
        y = fact(x-1);
        return x*y;
    }
}

int main(void) {
    int x = 0;
    x = fact(4);
    printf("Factorial 4 is %d", x);
    return 0;
}
```

		...
$SF_{main} -$	0xbfffffc : 0xbffff8 : 0xbffff4 :	0xb7e31a83 @ret0 0x00000000 %ebp0 0x00000000
$SF_{fact(4)} -$	0xbffffd0 : 0xbffffcc : 0xbffffc8 : 0xbffffbc :	0x00000004 0x08048474 @retm %ebp_m 0x00000001
$SF_{fact(3)} -$	0xbffffa0 : 0xbffff9c : 0xbffff98 : 0xbffff8c :	0x00000003 0x08048449 @ret4 0xbffffc8 %ebp4 0x00000001
$SF_{fact(2)} -$	0xbffff70 : 0xbffff6c : 0xbffff68 : 0xbffff5c :	0x00000002 0x08048449 @ret3 0xbffff98 %ebp3 0x00000001
$SF_{fact(1)} -$	0xbffff40 : 0xbffff3c : 0xbffff38 : 0xbffff2c :	0x00000001 0x08048449 @ret2 0xbffff68 %ebp2 0x00000001

x86 runtime memory

```
int fact(int x) {
    int y = 1;
    if (x == 0)
        return y;
    if (x > 0) {
        y = fact(x-1);
        return x*y;
    }
}

int main(void) {
    int x = 0;
    x = fact(4);
    printf("Factorial 4 is %d", x);
    return 0;
}
```

		...
$SF_{main} -$	0xbfffffc : 0xbffff8 : 0xbffff4 :	0xb7e31a83 @ret0 0x00000000 %ebp0 0x00000000
$SF_{fact(4)} -$	0xbffffd0 : 0xbffffcc : 0xbffffc8 : 0xbffffbc :	0x00000004 0x08048474 @retm 0xbffff8 %ebp_m 0x00000001
$SF_{fact(3)} -$	0xbffffa0 : 0xbffff9c : 0xbffff98 : 0xbffff8c :	0x00000003 0x08048449 @ret4 0xbffffc8 %ebp4 0x00000001
$SF_{fact(2)} -$	0xbffff70 : 0xbffff6c : 0xbffff68 : 0xbffff5c :	0x00000002 0x08048449 @ret3 0xbffff98 %ebp3 0x00000001
$SF_{fact(1)} -$	0xbffff40 : 0xbffff3c : 0xbffff38 : 0xbffff2c :	0x00000001 0x08048449 @ret2 0xbffff68 %ebp2 0x00000001
$SF_{fact(0)} -$		

x86 runtime memory


```
int fact(int x) {
    int y = 1;
    if (x == 0)
        return y;
    if (x > 0) {
        y = fact(x-1);
        return x*y;
    }
}

int main(void) {
    int x = 0;
    x = fact(4);
    printf("Factorial 4 is %d", x);
    return 0;
}
```

		...
$SF_{main} -$	0xbffeffc : 0xbffeff8 : 0xbffeff4 :	0xb7e31a83 @ret ₀ 0x00000000 %ebp ₀ 0x00000000
$SF_{fact(4)} -$	0xbffefd0 : 0xbffefcc : 0xbffefc8 : 0xbffefbc :	0x00000004 0x08048474 @ret _m 0xbffeff8 %ebp _m 0x00000001
$SF_{fact(3)} -$	0xbffefa0 : 0xbffef9c : 0xbffef98 : 0xbffef8c :	0x00000003 0x08048449 @ret ₄ 0xbffefc8 %ebp ₄ 0x00000001
$SF_{fact(2)} -$	0xbffef70 : 0xbffef6c : 0xbffef68 : 0xbffef5c :	0x00000002 0x08048449 @ret ₃ 0xbffef98 %ebp ₃ 0x00000001
$SF_{fact(1)} -$	0xbffef40 : 0xbffef3c : 0xbffef38 : 0xbffef2c :	0x00000001 0x08048449 @ret ₂ 0xbffef68 %ebp ₂ 0x00000001
$SF_{fact(0)} -$	0xbffef10 : 0xbffef0c : 0xbffef08 : 0xbffefec :	0x00000000 0x08048449 @ret ₁ 0xbffef38 %ebp ₁ 0x00000001

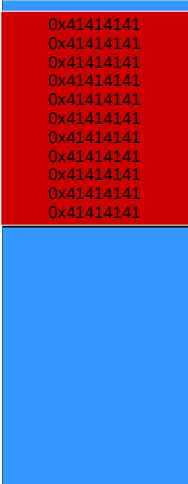
Buffer overflows

```
void function(char *str) {  
    char buffer[16];  
    strcpy(buffer,str);  
}  
void main() {  
    char large_string[256];  
    int i;  
    for( i = 0; i < 255; i++)  
        large_string[i] = 'A';  
    function(large_string);  
}
```



Buffer overflows

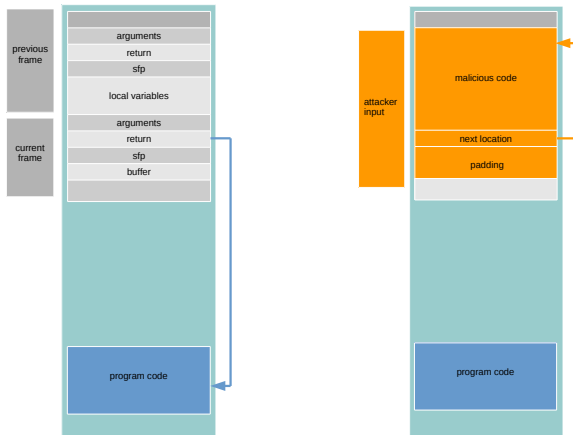
```
void function(char *str) {  
    char buffer[16];  
    strcpy(buffer,str);  
}  
void main() {  
    char large_string[256];  
    int i;  
    for( i = 0; i < 255; i++)  
        large_string[i] = 'A';  
    function(large_string);  
}
```



0x41414141
0x41414141
0x41414141
0x41414141
0x41414141
0x41414141
0x41414141
0x41414141
0x41414141
0x41414141

strcpy(src,dest) does not check that dest is bigger than src
The return address is now 0x41414141

Control hijacking



A buffer overflow can change the flow of execution of the program:

- ▶ load malicious code into memory
- ▶ make %eip point to it

Shellcode injection

Goal: "spawn a shell" - will give the attacker general access to the system

```
#include <stdio.h>
void main() {
    char *name[2];
    name[0] = "/bin/sh";
    name[1] = NULL;
    execve(name[0], name, NULL);
}
```

C code

```
"\x31\xc0"
"\x50"
"\x68" "//sh"
"\x68" "/bin"
"\x89\xe3"
"\x50"
...
```

Machine code

(part of attacker's input)

- ▶ must inject the machine code instructions (code ready to run)
- ▶ the code cannot contain any zero bytes (printf, gets, strcpy will stop copying)
- ▶ can't use the loader (we're injecting)

The return address

Challenge: find the address of the injected malicious code?

- ▶ If code accessible: we know how far is the overflowed variable from the saved %ebp
- ▶ If code not accessible: try different possibilities!
In a 32 bits memory space, there are 2^{32} possibilities
- ▶ NOP sled
 - ▶ guess approximate stack state when the function is called
 - ▶ insert many NOPs before Shell Code



Reference

Aleph One. Smashing The Stack For Fun And Profit.
<http://phrack.org/issues/49/14.html#article>